

How Do Language Models See Text?


A Primer on Tokenization

Many LLM Quirks Start with Tokenization

- **Letter counting:** a token can contain several letters, making character-level tasks less direct. For example, GPT-4 has incorrectly reported **two “r”s in “strawberry” instead of three**.
- **Arithmetic:** a number such as 12345 may split into 123 and 45, rather than individual digits.
- **Languages:** equivalent messages can require **different numbers of tokens**, affecting context usage.
- **Programming:** whitespace tokenization affects **how much code fits** into the context window.

Where the Tokenizer Fits

The Large Language Model (LLM) has its own training dataset, which is much larger than the dataset typically used to train the tokenizer.



The tokenizer will have its own training set of documents which are used to train the vocabulary using the Byte Pair Encoding (BPE) algorithm as a one-time procedure. Once we have the **vocabulary** and the **merges**, we can do both encoding and decoding.

A Brief History of Byte Pair Encoding

- **Byte Pair Encoding (BPE) originated as a compression algorithm.** It repeatedly replaces the most frequent pair of adjacent bytes with a single unused byte, producing a more compact representation (Gage, 1994).
- **Sennrich et al. adapted BPE for subword tokenization.** Their version merges characters and character sequences into reusable pieces, helping neural machine translation models represent **rare and unseen words** (Sennrich, Haddow & Birch, 2016).
- **GPT-2 used byte-level BPE for language modeling.** Starting from **256 possible byte values**, its tokenizer learns larger pieces and can encode any UTF-8 text without an unknown-character token (Radford et al., 2019).

How Does BPE Create New Tokens?

Example training text: hug hug hug pug pug

Step	“hug” × 3	“pug” × 2	New token added
Start	[h] [u] [g]	[p] [u] [g]	None
Merge u + g (5 occurrences)	[h] [ug]	[p] [ug]	[ug]
Merge h + ug (3 occurrences)	[hug]	[p] [ug]	[hug]
Merge p + ug (2 occurrences)	[hug]	[pug]	[pug]



Each separate block represents **one token**. We process each word separately.

<https://huggingface.co/learn/llm-course/en/chapter6/5>

How Training Data Shapes Tokenization

- **Different training data.** The tokenizer and LLM can use different datasets, with tokenizer training typically requiring much less text.
- **Frequent patterns get compact representations.** The training mix shapes which pieces enter the vocabulary.
- **Language balance matters.** English-heavy data can favour compact English representations over other languages.
- **Code and formatting matter too.** The training mix influences whitespace encoding and can favour formats such as YAML over JSON, or vice versa.
- **Vocabulary choices have lasting effects.** The vocabulary usually stays fixed during LLM training. Fewer tokens leave **more room in the context window**.

Numbers Split into Tokens of Variable Length

Tiktokenizer

gpt-4

System

You are a helpful assistant

×

User

Content

×

Add message

1234 + 5678 = 6912 # 123 and 567 are each single tokens.
112345 + 256789 = 369134 # 123 and 567 are NOT single tokens as
before, i.e. in the math statement, not in this very comment :).



Token count
59

1234 + 5678 = 6912 # 123 and 567 are each single tokens.
112345 + 256789 = 369134 # 123 and 567 are NOT single token
s as before, i.e. in the math statement, not in this very c
omment :).

4513, 19, 489, 220, 19282, 23, 284, 220, 25168, 17, 674, 22
0, 4513, 323, 220, 19282, 527, 1855, 3254, 11460, 627, 726
1, 12901, 489, 220, 4146, 16474, 284, 220, 19929, 9565, 67
4, 220, 4513, 323, 220, 19282, 527, 4276, 3254, 11460, 439,
1603, 11, 602, 1770, 13, 304, 279, 7033, 5224, 11, 539, 30
4, 420, 1633, 4068, 551, 570

<https://tiktokenizer.vercel.app/>

Number Tokenization in GPT-2 / GPT-3



X-axis: position within a row of 100 consecutive integers.

Y-axis: which block of 100 integers you are looking at. It increases **downwards**.

Moving **one cell right** advances by 1.

Moving **one row down** advances by 100.

Colour	Meaning	Example
Yellow: Unique	The entire number is one token	[number]
Red: 1-3	One digit, then three digits	2249 → [2][249]
Green: 2-2	Two digits, then two digits	2250 → [22][50]
Blue: 3-1	Three digits, then one digit	2251 → [225][1]

Check the full blog post at → <https://www.beren.io/2023-02-04-Integer-tokenization-is-insane/>

LLM Dyslexia

Tiktokenizer

gpt-4

System

You are a helpful assistant



User

Content



Add message

Token count

33

How many "r" letters does strawberry have?

strawberry vs strawberry --> note how a space in the front can change the tokenization and make counting hard.

[str][aw][berry] → 3 tokens

[strawberry] → 1 token



Same word, different IDs: the model sees token IDs, not letters.

For such a question, an agentic LLM can write and run a short piece of code that splits the word into letters and counts them.

It's important to develop an intuition about when to ask the LLM to do things via code and when to trust its outputs without additional tools.

Not All Languages Are Tokenized Equally

Tiktokenizer

gpt-4

System ▾

You are a helpful assistant

×

User ▾

Content

×

Add message

Hello

Bună ziua # A common Romanian expression used to greet people.

With this tokenizer, the Romanian greeting requires more tokens than the English greeting.

Since the context window is measured in tokens, reasoning in a language that requires more tokens for equivalent content can fill the model's context window more quickly.

Token count

63

Hello

Bună ziua # A common Romanian expression used to greet people.

With this tokenizer, the Romanian greeting requires more tokens than the English greeting.

Since the context window is measured in tokens, reasoning in a language that requires more tokens for equivalent content can fill the model's context window more quickly.

9906, 198, 33, 359, 6845, 64573, 4381, 220, 674, 362, 4279, 74697, 7645, 1511, 311, 41886, 1274, 382, 2409, 420, 47058, 11, 279, 74697, 43213, 7612, 810, 11460, 1109, 279, 6498, 43213, 382, 12834, 279, 2317, 3321, 374, 17303, 304, 11460, 11, 33811, 304, 264, 4221, 430, 7612, 810, 11460, 369, 13890, 2262, 649, 5266, 279, 1646, 753, 2317, 3321, 810, 6288, 13

<https://tiktokenizer.vercel.app/>

English dominates online text, so both the LLM and its tokenizer see more of it, and English ends up tokenized more compactly.

Tokenizing Source Code

Tiktokenizer

gpt2

```
for i in range(1, 11):  
    print(f"This is integer {i}")
```

Token count
24

```
for i in range(1, 11):  
    print(f"This is integer {i}")
```

Each space is encoded separately, resulting in longer token sequences

Tiktokenizer

System

User

Add message

```
for i in range(1, 11):  
    print(f"This is integer {i}")
```

gpt-4

Token count
19

```
for i in range(1, 11):  
    print(f"This is integer {i}")
```

Several spaces are merged together and treated as a single token.

In this example, the GPT-4 tokenizer groups consecutive spaces into **fewer tokens** than GPT-2's, allowing **more code to fit within the same token budget**.

UTF-8 Encoding

Code point ↔ UTF-8 conversion

First code point	Last code point	Byte 1	Byte 2	Byte 3	Byte 4
U+0000	U+007F	0yyyyzzzz			
U+0080	U+07FF	110xxxxyy	10yyzzzz		
U+0800	U+FFFF	1110wwww	10xxxxyy	10yyzzzz	
U+010000	U+10FFFF	11110uvv	10vvwwww	10xxxxyy	10yyzzzz

[<https://en.wikipedia.org/wiki/UTF-8>]

There are other encoding schemes like **UTF-16** and **UTF-32**. UTF-32, for example, is **fixed length**, whereas **UTF-8** is **variable-length**, which can result in shorter byte sequences. One of the major advantages of UTF-8 is that it is backwards compatible with **ASCII**, which is a much simpler encoding of text.

Check also → <https://www.reedbeta.com/blog/programmers-intro-to-unicode/>

Not All Byte Sequences Are Valid UTF-8

Example:

The character é (U+00E9) is encoded as two bytes: 11000011 10101001 → decimal values 195, 169

- 110 indicates the **start** of a two-byte sequence.
- 10 identifies the **continuation** byte.

By contrast, the byte 128 (10000000 in binary) cannot be decoded on its own: its 10 prefix marks it as a **continuation** byte, so a preceding **start** byte is required.

Hence, when you see lots of Unicode Replacement Characters like , the LLM post-processing pipeline cannot decode the bytes correctly for the browser to render. We join the bytes before decoding. If invalid byte sequences remain, `errors="replace"` displays  in their place:

```
def decode(token_ids):  
    # Join the bytes first: a UTF-8 character may span several tokens.  
    tokens = b"".join(vocab[token_id] for token_id in token_ids)  
    text = tokens.decode("utf-8", errors="replace") # In case of decoding issues use Unicode Replacement Characters  
    return text
```



Preventing Merges with Regex Pre-tokenization

Some merges should never take place. This can be enforced via regular expressions.

```
gpt-2 / src / encoder.py 
```

Code

Blame

117 lines (104 loc) · 4.14 KB



Raw

```
42     class Encoder:
43
44
45
46
47
48
49
50
51     # Should have added re.IGNORECASE so BPE merges can happen for capitalized versions of contractions
52     self.pat = re.compile(r"""'s|'t|'re|'ve|'m|'ll|'d| ?\p{L}+| ?\p{N}+| ?(?:\p{L}\p{N})+|\s+(?!\s)|\s+""")
53
54
55     def bpe(self, token):
56         if token in self.cache:
57             return self.cache[token]
58         word = tuple(token)
59         pairs = get_pairs(word)
60
61         if not pairs:
62             return token
63
```

The regex **splits text into smaller chunks before BPE is applied**. It recognizes:

- **English contractions:** 's, 't, 're, 've, 'm, 'll, 'd.
- **Groups of letters:** such as hello.
- **Groups of numbers:** such as 123.
- **Punctuation and symbols:** such as !!!.
- **Whitespace:** spaces, tabs and line breaks.

For example, "I'm buying 12 apples!" becomes: ["I", "m", " buying", " 12", " apples", "!"]

BPE then processes each chunk separately and cannot merge across these boundaries.

GPT-2 vs GPT-4 Pre-tokenization

tiktoken / tiktoken_ext / openai_public.py

Code Blame 162 lines (144 loc) · 5.48 KB

```
def cl100k_base():
    mergeable_ranks = load_tiktoken_bpe(
        "https://openaipublic.blob.core.windows.net/encodings/cl100k_base.tiktoken",
        expected_hash="223921b76ee99bde995b7ff738513eef100fb51d18c93597a113bcffe865b2a7",
    )
    special_tokens = {
        ENDOFTEXT: 100257,
        FIM_PREFIX: 100258,
        FIM_MIDDLE: 100259,
        FIM_SUFFIX: 100260,
        ENDOFPROMPT: 100276,
    }
    return {
        "name": "cl100k_base",
        "pat_str": r'(?i:[sdmt]|ll|ve|re)|[^\r\n\u{L}\u{N}]?+\u{L}++|\u{N}{1,3}+| ?[^\s\u{L}\u{N}]+[\r\n]*|\\s++$|\\s*[\r\n]|\s+(?!\s)|\s\"',
        "mergeable_ranks": mergeable_ranks,
        "special_tokens": special_tokens,
    }

def o200k_base():
    mergeable_ranks = load_tiktoken_bpe(
        "https://openaipublic.blob.core.windows.net/encodings/o200k_base.tiktoken",
        expected_hash="446a9538cb6c348e3516120d7c08b09f57c36495e2acfffe59a5bf8b0cfb1a2d",
    )
```

cl100k_base is the tokenizer used by the original GPT-4, GPT-4 Turbo, and GPT-3.5 Turbo

Example text: “_hello I'M buying 12345!\n” where, “\n” represents a new line / line break.

GPT-2: [“_”, “hello”, “I”, “”, “M”, “ buying”, “ 12345”, “!”, “\n”]

New regex: [“_hello”, “I”, “”M”, “ buying”, “ ”, “123”, “45”, “!\n”]

The new regex keeps “_hello”, “M”, and “!\n” together, but separates the space before the number and splits the digits into groups of up to three.

- English contractions are now case-insensitive: both 'm and 'M are recognized
- Numbers are split into groups of 1–3 digits: 12345 becomes "123" and "45", instead of one chunk
- Spaces before numbers are separated: " 42" becomes " " and "42"
- Words can include a preceding symbol: "_hello" can stay together, whereas GPT-2 separates "_ " and "hello"
- Punctuation can include following line breaks: "!\n" can stay together, whereas GPT-2 separates "!" and "\n"

GPT-4 Special Tokens

```
special_tokens = {  
    ENDOFTEXT: 100257,  
  
    FIM_PREFIX: 100258,  
  
    FIM_MIDDLE: 100259,  
  
    FIM_SUFFIX: 100260,  
  
    ENDOFPROMPT: 100276,  
}
```


Marks the end of a document or text sequence.

Ex: "Here is the text **before the gap**: *The cat*"

Ex: "Now **generate the missing text**."

Ex: "Here is the text **after the gap**: *on the mat.*"

Marks the end of the supplied prompt, after which the model generates its completion.


The cat sat on the mat.
fim_prefix fim_middle fim_suffix

Example:

- The model receives: <|fim_prefix|>The cat<|fim_suffix|> on the mat.<|fim_middle|>
- And then generates the missing middle: " sat"

Special Tokens Enable Tool Use

Tiktokenizer

cl100k_base

```
<|endoftext|> vs endoftext  
<|fim_prefix|> vs fim_prefix  
<|fim_middle|> vs fim_middle  
<|fim_suffix|> vs fim_suffix  
<|endofprompt|> vs endofprompt
```

Token count

27

```
<|endoftext|> vs endoftext  
<|fim_prefix|> vs fim_prefix  
<|fim_middle|> vs fim_middle  
<|fim_suffix|> vs fim_suffix  
<|endofprompt|> vs endofprompt
```

Note how the
special tokens are
a single block.

Special tokens are indivisible markers that structure messages and tool calls, e.g. `<message>` and `<call>` (illustrative names). Tool names, arguments, and content typically use **ordinary tokens**. The model learns the format during training, while external software executes tool calls.

100257, 6296, 842, 14450, 428, 198, 100258, 6296, 53187, 14
301, 198, 100259, 6296, 53187, 63680, 198, 100260, 6296, 53
187, 38251, 198, 100276, 6296, 842, 1073, 41681, 198

Search “*special_tokens*” → <https://github.com/openai/tiktoken/blob/main/src/lib.rs>

Structuring Conversations with Special Tokens

Tiktokenizer

gpt-4

System

You are a helpful assistant

×

User

What is BPE?

×

Add message

```
<|im_start|>system<|im_sep|>You are a helpful assistant<|im_end|>
<|im_start|>user<|im_sep|>What is BPE?<|im_end|>
<|im_start|>assistant<|im_sep|>
```

Role labels such as **system**, **user**, and **assistant** identify who is speaking. In this example, these labels are ordinary text enclosed by special tokens, so the model learns to interpret them as roles when they appear in that position.

Token count

21

```
<|im_start|>system<|im_sep|>You are a helpful assistant<|im_end|>
<|im_start|>user<|im_sep|>What is BPE?<|im_end|><|im_start|>assistant<|im_sep|>
```

```
100264, 9125, 100266, 2675, 527, 264, 11190, 18328, 100265,
100264, 882, 100266, 3923, 374, 426, 1777, 30, 100265, 100264, 78191, 100266
```

Special tokens mark the boundaries between messages and separate message headers from their content.

- `<|im_start|>` marks the beginning of a message.
- `<|im_sep|>` separates the message header from its content.
- `<|im_end|>` marks the end of a message.

The tiktoken Library



tiktoken

tiktoken is a fast [BPE](#) tokeniser for use with OpenAI's models.


```
import tiktoken
enc = tiktoken.get_encoding("o200k_base")
assert enc.decode(enc.encode("hello world")) == "hello world"

# To get the tokeniser corresponding to a specific model in the OpenAI API:
enc = tiktoken.encoding_for_model("gpt-4o")
```

The open source version of `tiktoken` can be installed from [PyPI](#):

```
pip install tiktoken
```

tiktoken is between 3-6x faster than a comparable open source tokeniser:



Performance measured on 1GB of text using the GPT-2 tokeniser, using `GPT2TokenizerFast` from `tokenizers==0.13.2`, `transformers==4.24.0` and `tiktoken==0.2.0`.

<https://github.com/openai/tiktoken>

Extending tiktoken

README MIT license



Extending tiktoken

You may wish to extend `tiktoken` to support new encodings. There are two ways to do this.

Create your `Encoding` object exactly the way you want and simply pass it around.

```
cl100k_base = tiktoken.get_encoding("cl100k_base")

# In production, load the arguments directly instead of accessing private attributes
# See openai_public.py for examples of arguments for specific encodings
enc = tiktoken.Encoding(
    # If you're changing the set of special tokens, make sure to use a different name
    # It should be clear from the name what behaviour to expect.
    name="cl100k_im",
    pat_str=cl100k_base._pat_str,
    mergeable_ranks=cl100k_base._mergeable_ranks,
    special_tokens={
        **cl100k_base._special_tokens,
        "<|im_start|>": 100264,
        "<|im_end|>": 100265,
    }
)
```


You can add more special tokens with a new ID afterwards.

Adding New Tokens to a Pretrained Model

1. Add the new token to the tokenizer.
2. Expand the input embedding matrix and output vocabulary projection.
3. Initialize the new parameters.
4. Fine-tune on examples containing the new token.

Adding special tokens can support adapting a **base model for chat** or introducing **tool-use formats** in chat and coding models. **Fine-tuning** teaches the model to interpret and generate these formats.

The SentencePiece Library



The screenshot shows the GitHub repository for SentencePiece. At the top, there are navigation links: README, Code of conduct, Contributing, Apache-2.0 license, and Security. The repository name "SentencePiece" is prominently displayed. Below it, a row of status badges indicates: CI for general build (passing), Build Wheels (passing), issues (3 open), python versions (3.9 | 3.10 | 3.11 | 3.12 | 3.13 | 3.14), and pyPI package (0.2.2). Another row shows: downloads (26M/month), contributions (welcome), License (Apache 2.0), SLSA (level 3), and docs (google.github.io/sentencepiece). The main text describes SentencePiece as a fast, lightweight, and unsupervised text tokenizer and detokenizer designed for neural network-based text generation systems. It mentions its ability to train directly from raw sentences and its implementation of subword units like BPE and unigram language models. A note states it is not an official Google product. The "Quick Start (Python)" section provides instructions to install the module via pip. The "Basic Example" section shows a code snippet for training a model and processing text.

SentencePiece

CI for general build **passing** Build Wheels **passing** issues **3 open** python **3.9 | 3.10 | 3.11 | 3.12 | 3.13 | 3.14** pyPI package **0.2.2**

downloads **26M/month** contributions **welcome** License **Apache 2.0** SLSA **level 3** docs **google.github.io/sentencepiece**

SentencePiece is a fast, lightweight, and unsupervised text tokenizer and detokenizer designed for neural network-based text generation systems (such as Large Language Models) where the vocabulary size is fixed prior to training.

It implements subword units—including Byte-Pair-Encoding (BPE) [Sennrich et al.] and the unigram language model [Kudo]—with the ability to train directly from raw sentences. By treating input text as a raw sequence of Unicode characters, SentencePiece enables a purely end-to-end, language-independent pipeline that completely eliminates the need for language-specific pre- or post-processing.

This is not an official Google product.

Quick Start (Python)

SentencePiece provides an easy-to-use Python module. Install it via `pip`:

```
pip install sentencepiece
```

Basic Example

Here is how to train a model, encode text into tokens/IDs, and decode them back to the original string:

```
import sentencepiece as spm

# 1. Train a model directly from a raw text file.
# (No pre-tokenization or language-specific preprocessing required!)
spm.SentencePieceTrainer.train(
    input='data/botchan.txt',
    model_prefix='m',
    vocab_size=1000
)

# 2. Load the trained model.
sp = spm.SentencePieceProcessor(model_file='m.model')
```

<https://github.com/google/sentencepiece>

- **SentencePiece** supports both **tokenizer training and inference**, whereas **tiktoken** focuses primarily on **inference with an already trained tokenizer**.
- **SentencePiece** supports multiple **algorithms**, including **BPE** and **Unigram**, whereas **tiktoken** focuses on **BPE**.
- **SentencePiece** is used by models such as **Llama 2**, **Mistral 7B**, and **Mixtral 8×7B**. (Llama 3 moved from SentencePiece to a tiktoken-based tokenizer)

Token Types in SentencePiece

1. **Special Tokens** → Reserved tokens with predefined functions, e.g. **<unk> for unknown tokens**, **<s>** for start of sentence, **</s>** for end of sentence, and optionally **<pad>** for padding.
2. **Byte Tokens** → Tokens representing individual bytes, used as a fallback for characters not covered by the normal vocabulary when **byte_fallback = True**
3. **Merge Tokens** → **Subword tokens** learned by repeatedly **merging** frequently occurring neighboring symbols during BPE training.
4. **Individual Codepoint Tokens** → Tokens representing individual Unicode characters retained according to the specified `character_coverage`. If a codepoint is too rare to fall within the cumulative character coverage specified by `character_coverage`, it is not added to the vocabulary, e.g. **character_coverage = 0.995**

Whitespace Handling in SentencePiece

Tiktokenizer

System ▾

You are a helpful assistant

×

User ▾

Content

×

Add message

dog is one of man's best friends.

the dog is one of man's best friends.

Token count

19

dog is one of man's best friends.

the dog is one of man's best friends.

18964, 374, 832, 315, 893, 596, 1888, 4885, 4286, 1820, 567
9, 374, 832, 315, 893, 596, 1888, 4885, 13

```
spm.SentencePieceTrainer.train(  
    input="token_training_text.txt",  
    model_prefix="model",  
    vocab_size=8000,  
    add_dummy_prefix=True  
)
```

If **add_dummy_prefix** is **True**, a whitespace marker (`_`) will be added at the start of the input. This treats the first word like a word preceded by a space, reducing the need to learn separate token variants for the same word with and without a preceding space.

tiktoken vs SentencePiece

	tiktoken	SentencePiece
Basic representation	UTF-8 bytes	Unicode text / pieces
Tokenization methods	BPE	BPE , Unigram, character-level, word-level
Pre-tokenization	regex constraints	SentencePiece normalization / whitespace scheme
Training	primarily prebuilt encodings; educational trainer exists	full tokenizer-training toolkit
Whitespace	regex/BPE-dependent	explicit “_” representation
Unknown characters	byte-based representation	“<unk>” or optional byte fallback

Takeaways

- **Tokens are the model's units of text.** The tokenizer converts text into a sequence of numerical IDs.
- **BPE builds tokens by merging frequent pairs.** Its training data shapes which patterns get compact representations.
- **Fewer tokens mean more room** in the context window. Different languages and formats use that space differently.
- **Token boundaries can be surprising.** Spaces affect the split, numbers break into chunks, and one character can span multiple tokens when its UTF-8 bytes fall into different tokens.
- **Special tokens provide structure.** The model learns the format during training, while external software executes tool calls.

References

Papers

- Gage (1994), *A New Algorithm for Data Compression* <https://jacobfilipp.com/DrDobbs/articles/CUJ/1994/9402/gage/gage.htm> [Slide 4, Speaker Notes Section]
- Sennrich, Haddow & Birch (2016). *Neural Machine Translation of Rare Words with Subword Units*. <https://arxiv.org/abs/1508.07909> [Slide 4, Speaker Notes Section]
- Radford et al. (2019). *Language Models are Unsupervised Multitask Learners* (GPT-2). https://cdn.openai.com/better-language-models/language_models_are_unsupervised_multitask_learners.pdf [Slide 4, Speaker Notes Section]
- Bavarian et al. (2022). *Efficient Training of Language Models to Fill in the Middle*. <https://arxiv.org/abs/2207.14255> [Slide 16, Speaker Notes Section]

Video

- Andrej Karpathy. *Let's build the GPT Tokenizer*. <https://www.youtube.com/watch?v=zduSFxRajkE> [Slide 23, Speaker Notes Section]

Code and libraries

- OpenAI GPT-2 encoder: <https://github.com/openai/gpt-2/blob/master/src/encoder.py> [Slide 14]
- tiktoken: <https://github.com/openai/tiktoken> [Slide 19]
- tiktoken encodings (cl100k_base regex): https://github.com/openai/tiktoken/blob/main/tiktoken_ext/openai_public.py [Slides 15, 16]
- tiktoken special tokens (lib.rs): <https://github.com/openai/tiktoken/blob/main/src/lib.rs> [Slide 17]
- Extending tiktoken: <https://github.com/openai/tiktoken#extending-tiktoken> [Slide 20]
- SentencePiece: <https://github.com/google/sentencepiece> [Slides 22, 23, 24, 25]
- Jupyter notebook: https://github.com/ducsp/UnderstandingTokenization/blob/main/BPE_Tokenizer_Notebook.ipynb [Slide 13]

Tools, courses and articles

- Tiktokenizer: <https://tiktokenizer.vercel.app/> [Slides 7, 9, 10, 11, 17, 18, 24]
- Hugging Face LLM Course, BPE chapter: <https://huggingface.co/learn/llm-course/en/chapter6/5> [Slide 5]
- Beren Millidge. *Integer tokenization is insane*. <https://www.beren.io/2023-02-04-Integer-tokenization-is-insane/> [Slide 8]
- Nathan Reed. *A Programmer's Introduction to Unicode*. <https://www.reedbeta.com/blog/programmers-intro-to-unicode/> [Slide 12]
- Wikipedia: UTF-8. <https://en.wikipedia.org/wiki/UTF-8> [Slide 12]